Inverting knowledge graphs back to raw data

How can we leverage the rules we use to construct knowledge graphs to do the inverse?

Tijs VAN KAMPEN

Promotor: Prof dr. ir. Anastasia Dimou

Master thesis submitted to obtain the degree of Master of Science in the engineering technology: Electronics-ICT

academic year 2023 - 2024

Abstract

Knowledge graphs are gaining traction nowadays and more and more companies use them, such as Amazon, Bosch, IKEA, Facebook, Google, LinkedIn, SIEMENS, Zalando, etc. Most knowledge graphs are nowadays constructed from other heterogeneous data sources, such as tables in relational databases, data in XML files, or JSON format derived from a Web API. While the construction of knowledge graphs from heterogeneous data has been thoroughly investigated so far, the inverse, namely constructing raw data from knowledge graphs hasn't been explored in depth yet.

In this thesis, we propose a method to invert knowledge graphs back to raw data. We will use the same rules used to construct the knowledge graph to do the inverse. Our method is split into two parts. First, a template is created by inverting the value retrieval of the mappings. For each supported source file this requires a tailored implementation. Secondly, the data is retrieved from the knowledge graph by using the mappings to generate a query for each iterator in the mappings. In the end, both these parts are combined, putting the data into the template.

The source code can be found on Github at https://github.com/TijsVK/Knowledge-graphs-inversion.

For this thesis, Github Copilot was used for assistance with the coding of the implementation, and rewording of sentences in the paper for readability.

Contents

Chapter 1

Introduction

The earliest academic definition of a knowledge graph can be found in a 1974 article as

> A mathematical structure with vertices as knowledge units connected by edges that represent the prerequisite relation (Marchi and Miguel, 1974; Bergman, 2019)

The idea of expressing knowledge in a graph structure predates even this definition, with the concept of semantic networks (Richens, 1956). However, the term knowledge graph only became well-known after Google announced they were using a knowledge graph to enhance their search engine in 2012 (Singhal, 2012). Knowledge graphs are used to make search engines, chatbots, question-answering systems, etc more intelligent by injecting knowledge into them (Ji et al., 2022).

A knowledge graph consists of many connected nodes, where each node is either an entity or a literal. These nodes are connected by edges, where each edge defines a relation between two nodes. Resource Description Framework (RDF) (Manola and Miller, 2004) is a framework often used to represent knowledge graphs, it uses subject-predicate-object triples to represent the nodes and their edges. Every subject node is either an URI or a blank node, while the object can be a literal value. The edges are URIs. This triple: `http://example.com/John_Doe` `http://schema.org/givenName "John"` . would represent the fact that the entity `John Doe` has the first name `John`. Often the predicates are chosen from an ontology/vocabulary, such as schema.org or FOAF. This allows for more interoperability between knowledge graphs, as the same predicates are used to represent the same concepts.

These knowledge graphs are constructed by extracting information from various sources, both unstructured sources such as text (using natural language processing) and (semi-)structured sources such as databases, CSV, XML, JSON, RDF (using mapping languages). Many mapping languages exist, differing in the way of defining the rules and the target source file format. Some mapping languages use the turtle syntax, while others provide their custom syntax, and others repurpose existing languages like SPARQL or ShEx (Van Assche et al., 2023). Some languages are specific to a single source format, such as R2RML(turtle format) (Das et al., 2012) for relational databases, XS-PARQL(XQuery format) (Bischof et al., 2012) for XML. Others can process multiple formats, such as RML (turtle) (Dimou et al., 2014), D-REPR (YAML) (Vu et al., 2019), xR2RML (turtle) (Michel

et al., 2015), etc. These can handle mapping from multiple sources in different formats.

To achieve the mapping of data these mapping languages use a declarative approach where the user specifies rules describing the desired output knowledge graph, the mapping rules. The implementation then takes care of the logic and transformations behind the mapping. Two ways of mapping exist, materialization and virtualization. Materialization constructs the knowledge graph as a file, which can be loaded into a triple store. Virtualization does not generate the knowledge graph as a file, but instead exposes a virtual knowledge graph, which can be queried as if it were a real knowledge graph. (Calvanese et al., 2017).

Creating these mapping rules is often done by hand. Some tools ease the creation process of these mappings, like RMLEditor (Heyvaert et al., 2018b) which exposes a visual editor, and YARRRML (Heyvaert et al., 2018a) which allows users to create rules in the user-friendly YAML which are then compiled to RML rules.

Retrieving data from a knowledge graph, for consumption by other programs, can be done by querying the knowledge graph using SPARQL (Seaborne and Prud'hommeaux, 2008). Using a select query data can be retrieved in a tabular format. A construct query can be used to retrieve the data in RDF format.

A knowledge graph cannot be converted back to the original data format using the same rules we created it with. As a result any changes we make to the knowledge graph are hard to propagate back. We can not update, expand, or improve the original data using e.g. knowledge graph refining nor can we apply changes to a virtual knowledge graph to change the original data.

In this work, we seek to answer the question: *How can we extend an existing mapping language like RML or create a new system to construct raw data from knowledge graphs?* We choose to extend the Morph-KGC implementation (Arenas-Guerrero et al., 2022) of RML (Dimou et al., 2014) as RML's end-to-end (from file to knowledge graph) characteristics make it a good candidate for this task. To answer the main research question we need to answer the following sub-questions:

RQ1 How can we construct the schema of the original data from the mapping rules?

  – We will study each type of source format, as each format has its challenges.

RQ2 How can we populate the schema with data from the knowledge graph?

  – We will study how we can best retrieve the data.

1.1 Thesis outline

The aim of this thesis is to explore the possibility of inverting knowledge graphs back to their original data format using RML mapping rules, we choose RML as it is well-adopted and has many implementations. To achieve this we will take a closer look at the technologies used like RDF, SPARQL, and RML in chapter 2. In chapter 3 a closer look will be taken at our implementation of the inversion algorithm. We will look at the algorithm itself, and the implementation details. In chapter 4 an evaluation of our implementation using various benchmarks will be done. For basic testing, we use a subset of the RML test cases, which are designed to test the conformance of tools to the RML

specification. For more advanced testing we will use various benchmarks simulating real-life use cases like LUBM4OBDA and GTFS-Madrid-Bench. Finally in chapter 5 we will conclude this thesis, and look at possible future work.

Chapter 2

Related work

In this chapter, we discuss the various technologies related to this thesis. We begin by discussing the semantic web and build from there to technologies used within its ecosystem like RDF, SPARQL, and mapping languages. We finish by discussing the current state of the art in updating or creating the original data source from a knowledge graph.

2.1 Semantic Web

Tim Berneers-Lee envisioned a version of the web that would also be understandable by machines, and thus the semantic web was born. It is not designed as a separate entity to the web, but instead as an extension, designed not for use by humans but rather computer 'agents'. It is designed with existing technologies like XML, URIs and RDF. Even ontologies, a key component of the semantic web, are not a new concept but rather co-opted from the field of philosophy. (Berners-Lee et al., 2001)

2.2 RDF

Resource Description Framework (RDF) was originally designed as a data model for metadata but has since been extended to be a general-purpose framework for graph data. RDF is a directed graph, where the nodes represent entities, and the edges represent relations between these entities. This graph is built up from triples, which connect a subject and an object using a predicate as shown in figure 2.1.

The subject must always be an entity, which can either be represented by an URI or be a blank node. The predicate must be an URI, and the object can be either an URI, a blank node, or a literal. (Manola and Miller, 2004)

A URI is a unique identifier for a resource on the web. Unlike a normal URL, it does not have to point to a network location, but can also be used to identify a person, a location, a concept, etc. (Manola and Miller, 2004). In RDF the URI is purely used for identifying resources. As such, unlike in HTML

Figure 2.1: An RDF triple

Subject	Predicate	Object
<http://example.com/De_Nayer>	<https://schema.org/location>	<http://example.com/Mechelen>.
<http://example.com/r0785695>	<https://schema.org/givenName>	"Tijs" .

Figure 2.2: Example of RDF triples

where certain conventions are expected, there are no conventions for URIs in RDF. An example of this can be seen in figure 2.2. In this figure the example subjects share the same domain but this does not imply that they are closely related, or even related at all. URIs are extended to IRIs to allow for a wider range of characters. Except for allowing Unicode characters, IRIs are identical to URIs so little distinction is made between the two in this thesis.

A blank node is a node that is not identified by a URI. It is used to represent an anonymous resource that can't be or has no reason to be uniquely identified. For example, the address of student r0785695 in figure 2.3 is only pertinent to the student and thus does not need to be uniquely identified. A blank node is serialized as `_:name`, where name is a unique identifier for the blank node. This identifier is only unique within the document, and thus can't be used to refer to the blank node from outside the document. (Manola and Miller, 2004)

A literal is a value, e.g. a string, integer, or date. This value can be typed, e.g. a string can be typed as a date, or untyped. A string can also have a language tag, which is used to indicate the language of the literal.

RDF is only a framework, and as such does not define any serialization syntax. There are however a few common serialization standards for example RDF/XML, Turtle, N-Triples, and JSON-LD.

2.2.1 Turtle

Turtle, or Terse RDF Triple Language, is a human-readable serialization format for RDF. It is the most used serialization format for RDF, and is used in many tools and specifications. In its sim-

Subject	Predicate	Object
ex:student/r0785695	schema:address	_:addrr0785695
_:addrr0785695	schema:postalCode	"2800"^^xsd:integer
_:addrr0785695	schema:streetAddress	"Gentsesteenweg XXXX"@nl

Figure 2.3: Example of a blank node and prefix notation

plest form turtle consists of triple statements, sequences of subject-predicate-object separated by spaces and terminated by a dot. An example of this can be seen in listing 2.1. This is very verbose, but Turtle offers many features to make it more concise. Below is a list of some of these features and, if possible, how they can be used to make the example more concise.

- **Prefix notation** allows us to shorten URIs by defining a prefix.
  - Using `@prefix schema: <https://schema.org/>` allows us to shorten `https://schema.org/Person` to `schema:Person`
- **Base prefix** allows us to shorten URIs by defining a base URI.
  - Using `@base <http://example.com/>` allows us to shorten `http://example.com/r0785695` to `r0785695`
- **Predicate lists** allow us to shorten multiple triples with the same subject to a list of predicates.
  - Our example only has two subjects, we can split their predicates with `;` instead of repeating the subject.
- **Object lists** allow us to shorten multiple triples with the same subject and predicate to a list of objects.
  - `r0785695` is both a Person and a Student, so we can split the objects with `,` instead of repeating the subject and predicate.
- **Literals** allow identifying values, e.g. strings, integers, dates, etc. with a datatype or language tag.
- **Blank nodes** allow us to define anonymous resources by using the `_:` prefix.
  - The address of `r0785695` is only relevant to `r0785695`, so we can define it as a blank node instead of using a URI, this shortens `<http://example.com/addrr0785695>` to `_:addrr0785695`. While not exactly the same, functionally it is equivalent as we do not expect addresses to be addressed outside of the context of a person.
- **Unlabeled blank nodes** allow us to define anonymous resources without explicitly defining an identifier by using the [] notation instead of `_:name`.
  - As we do not need to refer to `_:addrr0785695` from outside `r0785695`, we can use an unlabeled blank node and include it in `r0785695` instead of a labeled blank node.
- **Collections** allow us to define a list of nodes by using the () notation.

The example in listing 2.1 can be rewritten using these features, as shown in listing 2.2.

```
<http://example.com/r0785695> <http://www.w3.org/1999/02/22-rdf-syntax-ns#
    type> <http://schema.org/Person> .
<http://example.com/r0785695> <http://www.w3.org/1999/02/22-rdf-syntax-ns#
    type> <http://schema.org/Student> .
<http://example.com/r0785695> <http://schema.org/givenName> "Tijs" .
<http://example.com/r0785695> <http://schema.org/familyName> "Van Kampen" .
<http://example.com/r0785695> <http://schema.org/address> <http://example.
    com/addrr0785695> .
<http://example.com/addrr0785695> <http://schema.org/postalCode> "2800"^^<
    http://www.w3.org/2001/XMLSchema#integer> .
<http://example.com/addrr0785695> <http://www.w3.org/1999/02/22-rdf-syntax-
```

```
    ns#type> <http ://schema.org/PostalAddress> .
<http ://example.com/addrr0785695> <http ://schema.org/streetAddress> "
    Gentsesteenweg XXXX"@nl .
<http ://example.com/addrr0785695> <http ://schema.org/addressCountry> "
    Belgium" .
```

Listing 2.1: Basic naive turtle document

```
@prefix schema : <https ://schema.org/> .
@prefix xsd : <http ://www.w3.org/2001/XMLSchema#> .
@base <http ://example.com/> .
<r0785695> a schema:Person , schema:Student ;
    schema:givenName "Tijs" ;
    schema:familyName "Van Kampen" ;
    schema:address [
        a schema:PostalAddress ;
        schema:postalCode "2800"^^xsd:integer ;
        schema:streetAddress "Gentsesteenweg XXXX"@nl ;
        schema:addressCountry "Belgium"
    ] .
```

Listing 2.2: Basic turtle document using turtle features

2.3   SPARQL

SPARQL Protocol And RDF Query Language (SPARQL) is the W3C standard query language for
RDF. It is the main way to query RDF data and shows many similarities to SQL. SPARQL queries
mostly consist of a pattern of triples, which are matched against the RDF graph, a basic example
can be found in listing 2.3. Querying is very feature-rich, with support for aggregation, subqueries,
negation, regex, string manipulation, etc. It also supports different return types, federated queries,
entailment, etc (Harris and Seaborne, 2013). Aside from the query protocol it also defines the graph
store protocol, which can be used to manipulate graph databases directly (Aranda et al., 2013).

```
PREFIX schema : <https ://schema.org/>
SELECT ?name ?address
WHERE {
    ?student  schema:givenName ?name .
    ?student  schema:address ?address .
}
```

Listing 2.3: Example of a basic SPARQL SELECT query

2.4   Mapping languages

Mapping languages are used to define a mapping between a source and a target. The target in the context of linked data is of course RDF, with the source being any structured data source. Some mapping languages exist for a single source, e.g. Relational Database to RDF Mapping Language (R2RML) for relational databases, a query language combining XQuery and SPARQL (XSPARQL) for XML, etc. Others are more general purpose, e.g. RDF Mapping Language (RML) and Data to RDF Mapping Language (D2RML). We will discuss both R2RML and RML in more detail, as one extends the other. R2RML is significant as it is the most widely used mapping language, as it supports virtualization *ontop* of databases.RML extends R2RML with more features and source types making it one of the more feature-rich general mapping languages, as such it is the mapping language we will use in our implementation. Most RML implementations also support R2RML, as RML is a nearly superset of R2RML.

2.4.1   R2RML

Relational Database to RDF Mapping Language (R2RML) is a mapping language for mapping relational databases to RDF. As opposed to Direct Mapping (DM), which results in a direct mapping from the relational database to RDF without any changes to structure or naming, R2RML allows for more flexibility. R2RML mappings consist of zero or more TriplesMaps, which are used to map a table to RDF. A TriplesMap consists of a logical table, a subject map, and one or more predicate object maps (POMs). An example of an R2RML mapping can be seen in listing 2.4.

The logical table is used to define the table that is being mapped, with each row in the table being mapped to a subject and its corresponding POMs. It is possible to create a view of a table by using a SQL query, and then map this view. This allows for more complex mappings, e.g. mapping a join of two tables or a computed column.

Each of the SubjectMap, PredicateMap, ObjectMap, (and GraphMap) is a subclass of TermMap, which is a function that generates an RDF term. The map type can be constant, template, or column. The resulting term is then used as the subject, predicate, object, or graph of the triple. The termType of the map determines the type of the term, which can be IRI, blank node, or literal. If the termType is literal, optionally the datatype or language can be added. Following the RDF specification, not all combinations of termType and map are possible, this is shown in table 2.1. The object map has an additional subclass, a reference object map, in which we refer to another TriplesMap. Using a reference map we can map a foreign key to the subject of another TriplesMap, with a join condition. (Cyganiak et al., 2012)

TermType	Subject	Predicate	Object	Graph
IRI	✓	✓	✓	✓
Blank node	✓	✗	✓	✓
Literal	✗	✗	✓	✗

Table 2.1 Possible combinations of TermType and Map type

R2RML		RML	
Logical Table (database)	`rr:logicalTable`	Logical Source	`rml:logicalSource`
Table Name	`rr:tablename`	URI (pointing to the source)	`rml:source`
column	`rr:column`	reference	`rml:reference`
(SQL)	`rr:SQLQuery`	Reference Formulation	`rml:referenceFormulation`
per row iteration		defined iterator	`rml:iterator`

Table 2.2 Differences between R2RML and RML

A constant value is a fixed value, e.g. a URI or a string. A template is a string with placeholders, which are replaced by values from the logical row. A column is the value of a column in the logical row.

```
@prefix rr: <http://www.w3.org/ns/r2rml#>.
@prefix ex: <http://example.com/ns#>.

<#TriplesMap1>
    rr:logicalTable [ rr:tableName "EMP" ];
    rr:subjectMap [
        rr:template "http://data.example.com/employee/{EMPNO}";
        rr:class ex:Employee;
    ];
    rr:predicateObjectMap [
        rr:predicate ex:id;
        rr:objectMap [ rr:column "EMPNO";
                       rr:datatype xsd:positiveInteger ].
    ].
```

Listing 2.4: Example of an R2RML mapping

2.4.2 RML

RDF Mapping Language (RML) is a mapping language for mapping any (semi-)structured data source to RDF. It is a generalization of R2RML and as such supports all the features of R2RML. It extends R2RML by extending database-specific features to make them more general. The differences in usage can be seen in table 2.2.

RML uses the same structure as R2RML, with TriplesMaps consisting of a logical source, a subject map, and zero or more POMs. All the changes relate to the logical source. Whereas in R2RML the source is always a database, from which we select a table or view, in RML the source can be one of many different source types like XML, JSON, CSV, etc. Where in R2RML we simply iterate over the rows of a table, in RML we can have a source without an implicit iteration pattern, and as such we need to define an iterator.

2.5 State of the art

The state of the art in updating or creating the data source from a knowledge graph can be split in two categories, depending on the methodology used. The first methodology applies to virtualization, where the data is exposed as a virtual knowledge graph over the source data. The other methodology is for materialized knowledge graphs, where the knowledge graph is created as a file that can be loaded into a triple store. We will discuss the state of the art in both methodologies, concluding with the relevance of this work.

2.5.1 Virtualization

Virtualization is the process of exposing a virtual knowledge graph over the source data. This virtual knowledge graph can be queried as if it were a real knowledge graph. To achieve this, mappings are used to translate queries over the knowledge graph to queries over the source data.

This is most commonly used to expose a database as a virtual knowledge graph. This way an organization can expose a knowledge graph without having to completely transition to a new system. Most implementations of virtualization layers are read-only though, as the translation of SELECT queries is relatively easy but translating INSERT, DELETE or DELETE/INSERT (update) queries is much less straightforward, or even impossible in many cases. Though propagating changes trough the virtualization layer to the source data could be a big part of the linked data lifecycle, related work on this is scarce. In both "SPARQL Update queries over R2RML mapped data sources" (Unbehauen and Martin, 2017) and "Practical Update Management in Ontology-Based Data Access" (De Giacomo et al., 2017) the authors propose a similar way of handling updates. Compatible updates are propagated to the source data, while incompatible updates (updates which can not be translated back to the source because of conflicts) are held in an 'overflow' triple store. Further changes may make the incompatible updates compatible, at which point they too are propagated to the source data. Larger changes affecting the general structure of the data are not stored, but instead the mapping is updated to reflect the new structure. The handling of updates is shown in figure 2.4.

2.5.2 Materialization

Materialization is the process of creating a knowledge graph as a file that can be loaded into a triple store. The knowledge graph is then loaded into a dedicated triple store for querying. This method benefits from increased performance, at the cost of having the knowledge graph not in sync with the source data.

Materialization allows for a wider range of source formats, as the knowledge graph can be created from any structured data source. This is especially useful when the source data is not easily queryable, e.g. when the source data is a CSV, XML, or JSON file. For knowledge graphs created from structured data sources, use cases exist for propagating changes back to the original data source. This is not done using a direct update (as the source data is not directly connected to

Figure 2.4: Propagating changes in a virtualization layer (Unbehauen and Martin, 2017)

the knowledge graph), but by creating a new version of the source data. This process is called lowering. Below we discuss some of the state of the art in lowering.

2.5.2.1 XSPARQL

XSPARQL is a mapping language that allows for the transformation of XML to and back from RDF. It expands XQUERY with SPARQL-like syntax, structure and features. Its predecessors would do lifting by querying the source XML using XQuery to transform it into the XML serialization of RDF, while lowering was done using XSLT to do the inverse. Using its combined vocabulary XSPARQL simplifies lifting and lowering, using a single language for both and having the ability to target the turtle serialization (Polleres et al., 2009). An example lifting and lowering query can be found in listing 2.5 and listing 2.6 respectively.

```
declare namespace foaf="http://xmlns.com/foaf/0.1/";
declare namespace rdf="http://www.w3.org/1999/02/22-rdf-syntax-ns#";
let $persons := //*[@name or ../knows]
return

for $p in $persons
let $n := if( $p[@name] )
            then $p/@name else $p
let $id := count($p/preceding::*)
            +count($p/ancestor::*)
where
```

```
    not(exists($p/following::*[
        @name=$n or data(.)=$n]))
construct {
    _:b{$id} a foaf:Person;
            foaf:name {data($n)}.
    {
        for $k in $persons
        let $kn := if( $k[@name] )
                    then $k/@name else $k
        let $kid := count($k/preceding::*)
                    +count($k/ancestor::*)
        where
            $kn = data(//*[@name=$n]/knows) and
            not(exists($kn/../following::*[
                @name=$kn or data(.)=$kn]))
        construct {
        _:b{$id} foaf:knows _:b{$kid}.
        _:b{$kid} a foaf:Person.
        }
    }
}
```

Listing 2.5: Example of XSPARQL lifting

```
<relations>{
    for $Person $Name from <relations.rdf>
    where {$Person foaf:name $Name}
    order by $Name
    return
        <person name="{$Name}">{
            for $FName from <relations.rdf>
            where {
                $Person foaf:knows $Friend.
                $Person foaf:name $Name.
                $Friend foaf:name $Fname. }
            return
            <knows>{$FName}</knows>
        }</person>
}</relations>
```

Listing 2.6: Example of XSPARQL lowering

2.5.2.2   POSER: A Semantic Payload Lowering Service (Spieldenner., 2022)

POSER(Payload lOwering SERvice) is a service that lowers RDF to JSON. To achieve this a mapping is created in two parts: first the source patterns are defined from which to extract the data,

then the json structure is defined. The mapping is written in turtle, using a custom json ontology describing the json structure. A proof of concept implementation was made, but never got out of the prototype phase. It also only handles direct literal types, more complex composite values that are generated from templates are not supported. An example mapping is shown in listing 2.7.

```
@prefix ctd: <http://connectd.api/> .
@prefix onto: <http://ontodm.com/OntoDT#> .
@prefix iots: <http://iotschema.org/> .
@prefix json: <http://some.json.ontology/> .
@prefix xsd: <http://www.w3.org/2001/XMLSchema#> .
@prefix time: <https://www.w3.org/TR/2020/CR-owl-time-20200326/> .

# Which inputs to expect and to start mapping from
json:InputDataType {
    json:EntryPoint a iots:TimeSeries ;
        iots:providesTemperatureData iots:TemperatureData ;
        iots:providesTimeData iots:TimeData .

    iots:TemperatureData iots:numberDataType iots:Number .
    iots:TimeData   time:dateTime iots:Number .
}

#Semantic description of the json objects to be found in the expected API

json:ApiDescription {
    ctd:JsonModel json:hasRoot ctd:Node .

    ctd:TemperatureValue a json:Number ;
        json:key "value"^^xsd:string ;
        json:dataType iots:TemperatureData .

    ctd:TimeStamp a json:String ;
        json:key "timestamp"^^xsd:string ;
        json:dataType iots:TimeData .

    ctd:Node a json:Object;
        json:key "node"^^xsd:string ;
        json:value ctd:TimeStamp, ctd:TemperatureValue ;
        json:dataType iots:TimeSeries .

    ctd:Edges a json:Array        ;
        json:key "edges"^^xsd:string ;
        json:value ctd:Node .
}
```

Listing 2.7: Example of a POSER mapping

2.5.3   Relevance

As shown in this section the state of the art in updating the source in virtualization is pretty mature, only limited by fundamental limitations. It is however limited to databases. To work with other (semi-)structured data sources materialization is needed. For materialization the state of the art is much more limited. Though methods exist to lower RDF to other formats, each method is intimately linked to a source type. The mappings are also unidirectional, even XSPARQL which offers both lifting and lowering requires a separate mapping for each.

Our work expands RML, which can map from any structured data source to RDF, with the ability to lower the RDF back to the original source. We use a single mapping to do both lifting and lowering, as information on where to find the data during lifting can also be used to find where to put the data during lowering. This makes our work unique in the field of lowering RDF to other formats.

As the state of the art is the most limited for sources to which only materialization applies, this is where our work is most relevant. As such we choose to focus our implementation on this domain. However, the design and concept of our work can be applied to any structured data source, including database updates over a virtual graph.

Chapter 3

Implementation

In this chapter we present the design, limitations and implementation of the proposed system. The design and any of its components are not bound to any specific technology, as such they can be implemented using any programming language, framework or mapping language. Using a mapping outside its intended purpose has some limitations, we present these fundamental limitations in their own section. For the implementation we first present the general setup of the environment. We then present the implementation of the data retrieval. Lastly we present some implementations of template creation and filling.

3.1   System Design

The system we propose is a pipeline that takes as input a knowledge graph and set of mapping rules and outputs the source files from which it would have been constructed. It has the advantage of being very modular, consisting of separate data retrieval and template creation/filling modules. The data retrieval module takes a mapping file and a knowledge graph, and outputs a table with the necessary data. The templating module will take the same mapping file and the table outputted by the data retrieval module and output a source file. Depending on the mapping rules and config files different modules will be used. The graph source will decide the data retrieval strategy while the output file will determine the templating engine to be used. An overview of the design can be seen in figure 3.1.

We will implement the system in Python, using the Morph-KGC library to process the mapping rules. For the data retrieval module, we use RML to generate SPARQL queries to execute on a SPARQL endpoint (which we set up if necessary). This gives us a single table of data in CSV format for each source which we load into a Pandas DataFrame, which is then passed on to one of two templating modules we implement: one for tabular data (CSV) and one for nested/tree data (JSON). The structure for the templates is reconstructed from the RML mapping rules. Our specific implementation is shown in figure 3.2.
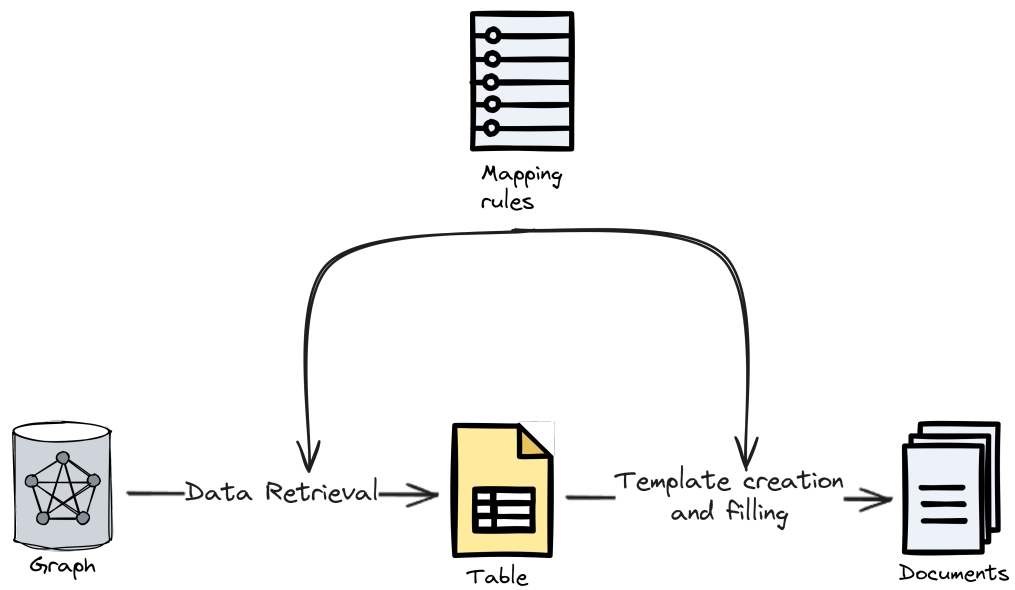
Figure 3.1: Design of the proposed system

Figure 3.2: Specific implementation of the design

3.2 System limitations

We observe three types of limitation inherent to the system. As they stem from the very basis of the system, some are impossible to overcome while others can be mitigated. The first two are caused by irreversible processing, transforming the source or data in such a way that it is impossible to reconstruct the original. The third is caused by the design of nested data structures, where the method of data access can limit the ability to reconstruct the source.

3.2.1 Irreversible source transformation

The first limitation is caused by the transformation of the source data before mapping. The most common irreversible transformation of a source in RML is a view over a database. While some views could be reversible, like a simple projection, many are not. To be able to fully recreate a static source it must be passed on in full to the mapper, without any aggregation, filtering, or other irreversible transformations. To update an existing source, having a unique identifying (set of) reference would suffice, for databases this would be the primary key(s). The mapping must also use all the data in the table. An example irreversible source mapping on a database can be found in listing 3.1. This limitation is not limited to database sources. Query languages working on data formats like JSON or XML also have the capability to filter or select data, making full reconstruction impossible. Some examples of this can be found in table 3.1.

```
<TriplesMap1>
    a rr:TriplesMap;

    rr:logicalTable [ rr:sqlQuery """
        SELECT "Name", COUNT("Sport") as SPORTCOUNT
        FROM "Student"
        """ ];

    rr:subjectMap [ rr:template
                  "http://example.com/resource/student_{\"Name\"}"; ];

    rr:predicateObjectMap
    [
        rr:predicate    foaf:name ;
        rr:objectMap    [ rr:column "\"Name\""; ];
    ];

    rr:predicateObjectMap
    [
            rr:predicate    ex:numSport ;
            rr:objectMap    [ rr:column "SPORTCOUNT"; ];
    ];
    .
```

Listing 3.1: Irreversible source mapping from R2RML Test cases

XPath	JSONPath	Description
//book[last()]	$..book[(@.length-1)] $..book[-1:]	the last book in order.
//book[position()<3]	$..book[0,1] $..book[:2]	the first two books
//book[isbn]	$..book[?(@.isbn)]	filter all books with isbn number
//book[price<10]	$..book[?(@.price<10)]	filter all books cheaper than 10

Table 3.1 Examples of irreversible source access in XPath and JSONPath

3.2.2 Irreversible data transformation

The second limitation is caused by the transformation of data during mapping. RML allows for the transformation of data using functions. These functions offer a lot of functionality, from string functions (concat, replace, etc.) to math to custom (self-defined) functions. The result of these functions can either be stored in the graph as a value or used to join data. This transformation is often irreversible, with information from the original data lost. Take for example `grel:toUpperCase`: this function transforms the string to an all uppercase string, here we lose information about the original case.

A special instance of this is the use of template maps. Template maps are a feature of most mapping languages, either explicitly (`d2rq:pattern`, `rr:template`) or implicitly through custom functions. These maps allow for the composition of multiple references into a single value using a template. This can be used to create both URIs (e.g. `http://example.com/resource/student_{"Name"}_{"First Name"}`) or literal values (e.g. `"{"Name"} {"First Name"}"`). For literal values this operation is theoretically irreversible as no guarantees can be made about the contents of the original data. Only by being sure the separator between the references consist of a pattern not present in the theoretical reconstructed source can we reliably split the value back into its original references. So to find out if we can reliably reconstruct the source we need the source, creating a chicken and egg problem. For URIs the operation is reversible if the separator consists of "reserved characters" defined by rfc3986 (Berners-Lee et al., 2005), these are as follows: `:/?#[]@!$&'()*+,;=`.

3.2.3 Unreconstructable data structures

The third limitation is caused by the design of nested data structures. To access the data multiple paths are valid to retrieve the data. The examples in table 3.1 also apply here, as they use 'recursive descent' to access the data. In the example any element with the name 'book' is taken from the JSON, no matter the depth. Alternatively a direct path could be written out like `$.store.book.title`. Another valid option for this particular dataset would be to use wildcard

operators like `$.*.*.title`. Only a direct path gives us enough information to reconstruct the source file using the mapping rules. A recursive descent gives no indications as to the depth of the results, and using wildcards gives no information about the naming of the elements. While a 'suggested' source file could still be reconstructed, and it would still work with the mapping rules, it would almost certainly not share the same structure as the original source.

3.3 Setup

We choose to make our implementation in Python as it both has an extensive set of libraries and is well suited for rapid development. We use Morph-KGC to process the mapping rules for its ease of use, being well made and its use of the pandas library to represent the mapping rules. As the graph source we choose a standalone SPARQL endpoint on a triple store. For working with graph (RDF) files, we first upload them to a triple store.

First the endpoint is set up, if necessary. If the source is a file, it is uploaded to the endpoint using the SPARQL 1.1 Graph Store HTTP Protocol. In our setup this is a locally running free version of GraphDB. The url of the repository is then wrapped using the SPARQLWrapper library and some settings are set.

Next we load the mapping rules from the mapping(RML) files using Morph-KGC's internal `retrieve_mappings` function. This function takes a config file (.ini format) as input, which specifies the location of the mapping files and various other settings which can be used to configure the behavior of the mapping. The mappings are returned as a pandas DataFrame which we enrich with various helper columns. An example of a single mapping rule (row in the DataFrame) can be found in listing 3.2. Marked in bold are the helper columns we add.

```
source_name: DataSource1
triples_map_id: #TM0
triples_map_type: http://w3id.org/rml/TriplesMap
logical_source_type: http://w3id.org/rml/source
logical_source_value: student.csv
iterator: nan
subject_map_type: http://w3id.org/rml/template
subject_map_value: http://example.com/{Name}
subject_references_template: http://example.com/([^\/]*)$
subject_references: ['Name']
subject_reference_count: 1
subject_termtype: http://w3id.org/rml/IRI
predicate_map_type: http://w3id.org/rml/constant
predicate_map_value: http://xmlns.com/foaf/0.1/name
predicate_references_template: None
predicate_references: []
predicate_reference_count: 0
object_map_type: http://w3id.org/rml/reference
object_map_value: Name
```

```
object_references_template:  None
object_references:  ['Name']
object_reference_count:  1
object_termtype:  http://w3id.org/rml/Literal
object_datatype:  nan
object_language:  nan
graph_map_type:  http://w3id.org/rml/constant
graph_map_value:  http://w3id.org/rml/defaultGraph
subject_join_conditions:  nan
object_join_conditions:  nan
source_type:  CSV
mapping_partition:  1-1-1-1
```

Listing 3.2: Example of a mapping rule in Morph-KGC

3.4 Retrieving the data

Retrieving the data is done by querying the SPARQL endpoint using an automatically generated query. All the required data from the source is retrieved at once. This does result in some degree of duplication for nested structures, but doing all processing and joining server-side is preferable. This sadly has a disadvantage of disjointed mappings currently resulting in a Cartesian product of the fragments. This is further discussed in subsection 3.4.2. The data is returned as a CSV-table which is then loaded into a pandas DataFrame for further processing. The only post processing done is the decoding of url-encoded strings where necessary.

3.4.1 Generating the queries

To generate the queries we first select the triple maps we want to use to generate the query. While constant values are always included, we can lighten the load on the server and in some cases increase reliability by reducing the amount and complexity of mapping rules we use. We design three operating modes:

- **Full**: All triple maps are used to generate the query. This is the most complete mode, where all present instance of each reference is checked. Template maps can cause issues when the template is properly not reversible.
- **Reduced**: All triple maps with a object map type of reference are used to generate the query. Only where necessary to retrieve all data template map types are used to try and avoid the issues with their reversibility.
- **Minimal**: Only a single datapoint is used for each reference, with a preference for reference object map types. This further reduces the load on the server and the complexity of the query.

We then generate the queries by translating the mapping rules into patterns. Each of the three map types, constant, reference, and template, generates a different pattern. For the constant map

type, we know that it will always be present with a constant value, so we can simply add it as a triple pattern like `?s $predicate $object_value`. Both reference and template map types will not generate a triple during materialization if any of the references are not present during the generation. We take this into account by allowing for blank fields using the `optional` keyword. This does however bring some problems with it, as we do want to check if multiple references to the same value contain the same value. The optional keyword would just fail silently if we try to assign to the same output variable, regardless of the value. As such we add some more logic using a combination of `BIND` and `FILTER` to check if the values of all occurrences of a reference are the same. Reference maps are also easy to work with as they directly translate back to the source.

As mentioned in subsection 3.2.2 template maps are not always reversible. We assume the user has followed the guidelines for template maps though, as we can only check the compliance to using reserved characters anyways. We first check if the graph value matches the template using regex. We then use string manipulation to split graph value into the different references. Here too we add some more logic using a combination of `BIND` and `FILTER` to check if the values of all occurrences of a reference are the same.

The algorithm for generating the queries can be found in algorithm 1. While abstracting some details, this follows the flow of the implementation. To give an example we will analyse and generate the query for the mapping rule in listing 3.3. The mapping rule consists of 4 triple maps: 1 constant, 2 reference, and 1 template.

3.4.1.1 SPARQL query generation example

Translating the constant map is simple, we add the triple pattern `?s <http://www.w3.org/1999/02/22-rdf-syn` `<http://xmlns.com/foaf/0.1/Person> .`.

The first reference map uses a value not used in an encoded format, as such we do not have to encode it. We ensure that each value of `?amount` is the same, to achieve this we first assign it to a temporary variable with `OPTIONAL{?s <http://example.com/owes> ?amount_temp}` and then try to bind it to `?amount` with `OPTIONAL{BIND(?amount_temp as ?amount)}`. We then filter to make sure that if the value is present it matches all occurances with: `FILTER(!BOUND(?amount)` `|| !BOUND(?amount_temp) || ?amount_temp = ?amount)`. If `?s <http://example.com/owes>` has no value `?amount_temp` will be unbound, making `?amound` unbound too, making the filter pass. If `?s <http://example.com/owes>` does have a value, but `?amount` is already bound, the binding will fail but the filter will check if the values match. Lastly if `?s <http://example.com/owes>` has a value and `?amount` is unbound, the binding will succeed and the filter will (redundantly) check if the values match.

The second reference will be very similar to the first, but as the value is used in an encoded format in the template map we encode it using `ENCODE_FOR_URI` before binding and filtering.

The template map is more complex, as we have to split the value into the different references. We start by checking if the string matches the format of the template using a regex, this regex string is created by replacing each place a value is inserted with match-all quantifier. We then iterate

over the references, each time binding the reference and creating the next segment of the string. The first segment is created by slicing the string after the base like: `BIND(STRAFTER(STR(?s), 'http://example.com/') as ?temp)`. We know that the part up to the first separator is the first reference, so we bind it to a temporary variable with `BIND(STRBEFORE(STR(?temp), ';') AS ?fname_temp2)`. Much like the reference maps we use `BIND`ing and `FILTER`ing to ensure that each value of the reference is the same. We also know that the value after the separator is the second reference, as there are no more separators after it. Analog to the first reference we bind and filter the value.

Algorithm 1 Generating the queries

Require: $mapping\_rules$ is a set of mapping rules

1: $query\_lines \leftarrow []$
2: **for all** $rule \in mapping\_rules$ **do**
3:     $object\_encoded = rule.is\_encoded()$
4:     **if** $rule.is\_constant()$ **then**
5:         $query\_lines.append(rule.to\_triple())$
6:     **else if** $rule.is\_reference()$ **then**
7:         $query\_lines.append(rule.to\_optional\_triple(encode = object\_encoded))$
8:     **else if** $rule.is\_template()$ **then**
9:         $query\_lines.append(test\_object\_regex(rule))$
10:        $remainder \leftarrow rule['object\_template']$
11:        **for all** $reference \in rule['object\_references']$ **do**
12:            $query\_lines.append(bind\_next\_partial(rule, reference))$
13:            $remainder \leftarrow next\_segment(remainder)$
14:            **if** $remainder = ""$ **then**
15:                $query\_lines.append(rule.bind\_last\_slice())$
16:            **else**
17:                $query\_lines.append(rule.bind\_next\_slice())$
18:            **end if**
19:        **end for**
20:     **end if**
21: **end for**
22: $query \leftarrow wrap\_query\_lines(query\_lines)$

```
@prefix rr: <http://www.w3.org/ns/r2rml#> .
@prefix foaf: <http://xmlns.com/foaf/0.1/> .
@prefix ex: <http://example.com/> .
@prefix xsd: <http://www.w3.org/2001/XMLSchema#> .
@prefix rml: <http://semweb.mmlab.be/ns/rml#> .
@prefix ql: <http://semweb.mmlab.be/ns/ql#> .
@base <http://example.com/base/> .

<TriplesMap1> a rr:TriplesMap;
```

```
rml:logicalSource [
    rml:source "ious.csv";
    rml:referenceFormulation ql:CSV
];

rr:subjectMap [
    rr:template "http://example.com/{fname};{lname}";
    rr:class foaf:Person;
];

rr:predicateObjectMap [
    rr:predicate ex:owes;
    rr:objectMap [ rml:reference "amount"; ]
].

rr:predicateObjectMap [
    rr:predicate ex:firstName;
    rr:objectMap [ rml:reference "fname"; ]
].
```

Listing 3.3: Example mapping rule using all three map types

```
SELECT ?amount ?lname_encoded ?fname_encoded WHERE {
    ?s <http://www.w3.org/1999/02/22-rdf-syntax-ns#type>
        <http://xmlns.com/foaf/0.1/Person> .

    OPTIONAL{?s <http://example.com/owes> ?amount_temp}
    OPTIONAL{BIND(?amount_temp as ?amount)}
    FILTER(!BOUND(?amount) ||
            !BOUND(?amount_temp) ||
            ?amount_temp = ?amount)

    OPTIONAL{?s <http://example.com/firstName> ?fname_temp1}
    OPTIONAL{BIND(ENCODE_FOR_URI(?fname_temp1) as ?fname_encoded)}
    FILTER(!BOUND(?fname_encoded) ||
            !BOUND(?fname_temp1) ||
            ENCODE_FOR_URI(?fname_temp1) = ?fname_encoded)

    FILTER(REGEX(STR(?s), 'http://example.com/([^\\/]*);([^\\/]*)'))
    {} OPTIONAL{BIND(STRAFTER(STR(?s), 'http://example.com/') as ?temp)}
    BIND(STRBEFORE(STR(?temp), ';') AS ?fname_temp2)
    {} OPTIONAL{BIND(?fname_temp2 as ?fname_encoded)}
    FILTER(!BOUND(?fname_encoded) ||
            !BOUND(?fname_temp2) ||
            ?fname_temp2 = ?fname_encoded)
    {} OPTIONAL{BIND(STRAFTER(STR(?temp), ';') as ?lname_temp)}
```

```
    {} OPTIONAL{BIND(?lname_temp as ?lname_encoded)}
    FILTER(!BOUND(?lname_encoded) ||
            !BOUND(?lname_temp) ||
            ?lname_temp = ?lname_encoded)
}
```

Listing 3.4: Generated query for the mapping rule in listing 3.3

3.4.2 Disjointed mappings

We generate a query to retrieve the whole source at once, this can lead to issues when subjects in the mapping has no path/link between them. The effect this creates is not unlike joining two tables in SQL without join conditions. For example, using the mapping listed in listing 3.5 we get the query in listing 3.6. When applied to the knowledge graph in listing 3.7 we get the result in listing 3.8 instead of the original source in listing 3.9. When converting the badly generated source back to the knowledge graph, we do get the same knowledge graph as the original as the duplicate data is ignored. The amount of returned rows is a cartesian product between the fragments.

This is solvable by analysing the mapping rules to look for disjointed mappings and split them into separate queries. At the templating stage we could then merge the results back together. This would solve the issue of duplicated fragments, but the unmapped connections that may have been present in the source file will be lost. An example for how the result would look can be found in listing 3.10.

```
<TriplesMap1> a rr:TriplesMap;
rml:logicalSource [
    rml:source "student_sport.csv";
    rml:referenceFormulation ql:CSV
];
rr:subjectMap [
    rr:template
        "http://example.com/{Student}";
    rr:class ex:Student
];
rr:predicateObjectMap [
    rr:predicate foaf:name ;
    rr:objectMap [
        rml:reference "Student"
    ]
].
```

```
<TriplesMap2> a rr:TriplesMap;
rml:logicalSource [
    rml:source "student_sport.csv";
    rml:referenceFormulation ql:CSV
];
rr:subjectMap [
    rr:template
        "http://example.com/{Sport}";
    rr:class ex:Sport
];
rr:predicateObjectMap [
    rr:predicate foaf:name ;
    rr:objectMap [
        rml:reference "Sport"
    ]
].
```

Listing 3.5: Bad join mapping

```
SELECT DISTINCT ?Student_name ?Sport
```

```
WHERE {
    ?s1 a ex:Student .
    optional{
        ?s1 foaf:name ?Student_name .
    }
    ?s2 a ex:Sport .
    optional{
        ?s2 foaf:name ?Sport .
    }
}
```

Listing 3.6: Bad join query (trimmed)

```
@prefix ex: <http://example.com/> .
@prefix foaf: <http://xmlns.com/foaf/0.1/> .

ex:Venus a ex:Student ;
    foaf:name "Venus" .
ex:Tom a ex:Student ;
    foaf:name "Tom" .
ex:Tennis a ex:Sport ;
    foaf:name "Tennis" .
ex:Football a ex:Sport ;
    foaf:name "Football" .
```

Listing 3.7: Bad join knowledge graph

```
Student, Sport
Venus, Tennis
Venus, Football
Tom, Tennis
Tom, Football
```

Listing 3.8: Bad join result

```
Student, Sport
Venus, Tennis
Tom, Football
```

Listing 3.9: Bad join original source

```
Student, Sport
Venus,
Tom,
, Tennis
, Football
```

Listing 3.10: Better bad join result

3.5 Contructing the schema

Constructing the schema is done by reversing the mapping rules' source. We retrace the path where the data came from, using the mapping rules and generate a template. We then take the data and apply it to this template. The resulting source file is then written away. Most RML mappers support various referenceFormulations, which are used to describe the structure and method of access of the source. We implemented the CSV source type for the PoC, and now also implement the JSON source type.

As each referenceFormulation or even source type has its reference syntax, we tailor the implementation to each type. We implement the CSV source type as an implementation of table-like sources, and the JSON source type as an implementation of nested/tree sources. Other source types can be implemented later using similar principles or further processing.

3.5.1 CSV

The CSV referenceFormulation not only applies to CSV files, but also for any two-dimensional table-like structure like TSV, databases, Excel, etc. We however only implement the CSV source type, implementing other CSV-like sources would be very similar or simply more specific without adding any new insights. The example TriplesMap in listing 3.11 results in the CSV template in listing 3.12.

```
<TriplesMap1> a rr:TriplesMap;

rml:logicalSource [
    rml:source "student.csv";
    rml:referenceFormulation ql:CSV
];

rr:subjectMap [
    rr:template "http://example.com/Student/{ID}/{Name}";
    rr:graph ex:PersonGraph ;
    rr:class foaf:Person
];

rr:predicateObjectMap [
    rr:predicate ex:id ;
    rr:objectMap [ rml:reference "ID" ]
];

rr:predicateObjectMap [
    rr:predicate foaf:name ;
    rr:objectMap [ rml:reference "Name" ]
].
```

Listing 3.11: Example mapping for a CSV file

```
ID ,Name
<ID>,<Name>
```

Listing 3.12: Example CSV template

To apply the CSV data to the template we have to take no further steps, after the templating-CSV-module received the data already in a table format. We can simply write the table to a CSV file. As we use pandas to load the data, we use the built-in `to_csv` function to achieve this.

3.5.2 JSON

We use the JSON referenceFormulation as an example of a nested source. We don't think implementing other nested sources like XML would bring many new challenges or insights. The example YARRRML mapping in listing 3.13 generates the JSON template in listing 3.14, we use YARRRML here to keep the example concise.

```
prefixes:
    ex: "http://example.com/"

mappings:
    program:
        sources:
            - ['data.json~jsonpath', '$.Programs[ * ]']
        s: http://example.com/program/$(Name)
        po:
            - [a, ex:Program]
            - [ex:name, $(ProgramName)]

    course:
        sources:
            - ['data\.json~jsonpath', '$.Programs[ * ].Courses[ * ]']
        s: http://example.com/course/$(Name)
        po:
            - [a, ex:Course]
            - [ex:name, $(CourseName)]
            - [ex:credits, $(Credits)]
            - p: ex:partOf
            o:
                - mapping: program
                condition:
                    function: equal
                    parameters:
                    - [str1, $(CourseName)]
                    - [str2, $(Courses.CourseName)]
```

Listing 3.13: Example YARRRML mapping for a JSON file

ProgramName	CourseName	Credits
Program1	Course1	3
Program1	Course2	4
Program1	Course3	2
Program2	Course7	3
Program2	Course32	4

Table 3.2 Example JSON table

```
{
    "Programs": [
        {
            "Courses": [
                {
                "CourseName": "$CourseName",
                "Credits": "$Credits"
                }
            ],
            "ProgramName": "$ProgramName"
        }
    ]
}
```

Listing 3.14: Example JSON template

The templating module receives the data in form of a table. Each row contains the values for a full path into the nested structure. An example of a table for the mapping in listing 3.13 can be found in table 3.2. In the JSON structure we define two types of nodes, Objects and Arrays. Objects contain key-value pairs, each value being either a primitive or another node. Arrays contain a list of components sharing the same form. To fill the template we pass the data to the root node of the template, from there it is recursively passed down to the children. At each array node rows are grouped and split by the primitives filled in the layers below the array node up to the next array node and then iteratively passed down to its child node.

For the simple example below the template will recursively pass on the data until it reaches the array in the "Programs" key. It will then check which primitives will be filled up to the next array node, in this case just the "ProgramName". The array node then groups the rows by the "ProgramName" column and then iteratively passes each group to its child node. The child node will fill the "ProgramName" key and then pass the data to the "Courses" keys value. The array node in the "Courses" key will then group the rows by the "CourseName" and "Credits" columns and then iteratively pass each group to its child node. Each node ensures that the returned string from its child node is properly formatted and then returns it to its parent. The iteration steps can be found in figure 3.3.

- Program passes full table to root node
- Root node passes full table to "Programs" key child node (Array)
- Programs array node finds the primitives to fill before the next array node, in this case "ProgramName"
- Programs array node groups the rows by the "ProgramName" column
- Programs array node passes group with "ProgramName" = "Program1" to its child node (3 rows)
- Program object fills the "ProgramName" key and passes the data to the "Courses" key
- Courses array node finds the primitives to fill before the next array node, in this case "CourseName" and "Credits"
- Courses array node groups the rows by the "CourseName" and "Credits" columns
- Courses array node passes group with "CourseName" = "Course1" and "Credits" = "3" to its child node (1 row)
- Course object fills the "CourseName" and "Credits" keys and returns its string to the Courses array node
- Courses array node passes group with "CourseName" = "Course2" and "Credits" = "4" to its child node (1 row)
- ...
- After filling all rows the Courses array node returns its string to the Program object node
- The Program object node returns its string to the Programs array node
- Programs array node passes group with "ProgramName" = "Program2" to its child node (2 rows)
- ...
- After filling all rows the Programs array node returns its string to the root node
- The root node returns the full string

Figure 3.3: Iteration steps for the JSON template

Chapter 4

Evaluation

In this chapter we will evaluate our implementation. This will be done by testing it again various datasets, comparing the expected results with the actual results. For each dataset we go over our testing methodology and its results.

4.1 RML test cases

The RML test cases (Heyvaert et al., 2019) are a set of test cases to evaluate the conformity of an RML processor. Though these test cases are not a perfect match, they offer expected outputs for certain inputs and mapping rules, making them a good starting point for testing our implementation. The test cases are designed with edge cases in mind, and while they are an appropriate set of test cases to test a mapper, using them to test inversion is stretching their purpose a bit. As they are made to check if the full specification is implemented, many tests have duplicate inputs and outputs with differently written mapping rules to test the lexer, which we didn't make ourselves. They also do not conform to the limitations discussed in section 3.2. As such, if we were to simply take these test cases and invert them, we will get many generated source files that do not match the originals. Instead we try generating the knowledge graph again from our generated source files and compare the results with the original knowledge graph. We test the CSV and JSON test cases, as those are the ones we implemented.

4.1.1 CSV test cases

For the CSV test cases 23 out of 32 tests pass, the full breakdown can be found in figure 4.2. Most of the failures are limitations of the mapping processor, which only future update to the Morph-KGC library or a change to a different processor could solve. Another failure is due to data being stored in a blank node identifier, but no guarantees are made about the blank node's identifier in the RDF specification. For specific triple stores, it might be possible to retrieve the data from the identifier, but this is not possible as a general approach. The other two failures are due to the data not being stored at the subject, rather only being used for a join condition. In this specific case the data could

be retrieved as the join condition is an 'equals' function and the data is stored in the joined subject.

Out of 32 tests:
- 23 tests pass
- 2 fail due to data being stored in a blank node (e.g. listing 4.1)
- 2 fail due to data not being directly stored at the subject, but being used for a join condition (e.g. listing 4.3)
- 5 fail because the mapping processor does not support them

Figure 4.1: Results of the CSV RML test cases

4.1.2 JSON test cases

For the JSON test cases 24 out of 34 tests pass, the full breakdown can be found in figure 4.2. The failures are similar to the CSV test cases, aside from an extra failure due to the mapping containing no references, crashing the templating engine (for CSV an empty file is generated without crashing)

Out of 34 tests:
- 24 tests pass
- 2 fail due to data being stored in a blank node (e.g. listing 4.1)
- 1 fails because the mapping contains no references, crashing the templating engine (e.g. listing 4.2)
- 2 fail due to data not being directly stored at the subject, but being used for a join condition (e.g. listing 4.3)
- 5 fail because the mapping processor does not support them

Figure 4.2: Results of the JSON RML test cases

```
<TriplesMap1> a rr:TriplesMap;
    rml:logicalSource [
        rml:source "student.csv";
        rml:referenceFormulation ql:CSV
    ];

    rr:subjectMap [
        rr:template "students{ID}";
        rr:termType rr:BlankNode
    ];

    rr:predicateObjectMap [
        rr:predicate foaf:name;
        rr:objectMap [ rml:reference "Name" ]
    ].
```

Listing 4.1: Example RML mapping with data being stored only in a blank node

```
<TriplesMap1> a rr:TriplesMap;
    rml:logicalSource [
        rml:source "student.json";
        rml:referenceFormulation ql:JSONPath;
        rml:iterator "$.students[ * ]"
    ];

    rr:subjectMap [
        rr:constant ex:BadStudent;
        rr:graphMap [ rr:constant <http://example.com/graph/student> ];
    ];

    rr:predicateObjectMap [
        rr:predicateMap [ rr:constant ex:description ];
        rr:objectMap [ rr:constant "Bad Student"; ]
    ].
```

Listing 4.2: Example RML mapping with no references

```
TriplesMap1:
    sources:
    – [student.csv˜csv]
    s: http://example.com/resource/student_$(ID)
    po:
    – [http://xmlns.com/foaf/0.1/name, $(Name)]
    – p: http://example.com/ontology/practises
        o:
        mapping: TriplesMap2
        condition:
            function: equal
            parameters:
            – [str1, $(Sport)]
            – [str2, $(ID)]
TriplesMap2:
    sources:
    – [sport.csv˜csv]
    s: http://example.com/resource/sport_$(ID)
    po:
    – [http://www.w3.org/2000/01/rdf–schema#label, $(Name)]
```

Listing 4.3: Example YARRRML mapping with data only being used in a join condition

4.2   LUBM4OBDA

The LUBM for Ontology Based Data Access (LUBM4OBDA) benchmark (Arenas-Guerrero et al., 2024) is an extension of the Lehigh University Benchmark (LUBM) benchmark (Guo et al., 2005). Instead of generating OWL data it generates sql data, which paired with R2RML and RML mappings can be used to test OBDA systems. We use the generated sql data to test the performance of our implementation on different scales. As we do not implement a database module to recreate databases from a graph we instead reconstruct the views over the database which are used in the mappings. Comparing the speed the inversion is done for different scales of the benchmark gives us an idea of how well our implementation scales.

4.2.1   Accuracy

We compare the generated source files with the views of the mapping. We find that 15 out 22 source files are successfully generated. The issues are caused by the mapping rules having duplicate subject-predicate-object maps generated by different sources without constants to differentiate them. An example conflict can be seen in Listing 4.4. As this is a flaw in the mapping rules, we can not mitigate this issue.

```
<#GraduateStudentAdvisor>
    rml:logicalSource [
        rml:source "graduatestudentadvisor.csv" ;
        rml:referenceFormulation ql:CSV ;
    ];
    rr:subjectMap [
        rr:template "http://www.department{dnr}.university{unr}.edu/{gname}";
    ];
    rr:predicateObjectMap [
        rr:predicate ub:advisor;
        rr:objectMap [ rr:template
                "http://www.department{dnr}.university{unr}.edu/{fname}" ];
    ].

<#UndergraduateStudentAdvisor>
    rml:logicalSource [
        rml:source "undergraduatestudentadvisor.csv" ;
        rml:referenceFormulation ql:CSV ;
    ];
    rr:subjectMap [
        rr:template "http://www.department{dnr}.university{unr}.edu/{ugname}";
    ];
    rr:predicateObjectMap [
        rr:predicate ub:advisor;
        rr:objectMap [ rr:template
                "http://www.department{dnr}.university{unr}.edu/{fname}" ];
```

```
    ].
```

Listing 4.4: Example of a duplicate mapping pattern in the LUBM4OBDA benchmark

4.2.2 Performance

We run our program on the LUBM4OBDA benchmark for different scales. The test is run on a machine with a Ryzen 7 7800x3D processor and 64GB of RAM. We use the free version of GraphDB as our triple store, this limits us to single threaded performance. The results can be found in Table 4.1. We find that the time it takes to invert the mappings scales linearly with the scale of the benchmark. The time spent within the program for data retrieval is minimal by design, leaving the majority of the computations to the triple store. As such, inverting the mappings is mostly dependent on the triple store's performance. The time required to convert the mappings to the query is minimal, even at the smallest scale it constitutes less than 0.2% of the total time. For CSV files, minimal conversion is needed to transform the data to the source files. Although this time scales linearly with the benchmark scale, it accounts for less than 0.1% of the total time.

Scale	Time
1	12.49s
10	127.61s
100	1379.26s
1000	13826.61s

Table 4.1 Performance on the LUBM4OBDA benchmark

4.3   GTFS-Madrid-Bench

The GTFS-Madrid-Bench (Priyatna et al., 2023) benchmark is a benchmark for benchmark evaluating declarative KG construction engines. The data sources are based on the General Transit Feed Specification (GTFS) data files of the subway network of Madrid. This data can be transformed into several formats such as CSV, JSON, SQL and XML. A scaling factor can also be applied to the data, allowing for different sizes of the benchmark. We use the CSV and JSON data sources to test our implementation on various scales. This is a good benchmark to test the performance of our JSON templating engine, as we have a baseline for the duration of the data retrieval with the CSV files.

4.3.1   Accuracy

Comparing the generated to the original source files, we find that 7 out of 10 source files fully match for CSV, and 5 out of 10 source files match for JSON. The extra mismatches in JSON are due to the formatting of numbers, which is not guaranteed to be the same. In the original data the number

are formatted as integers even though they are converted to doubles in the RDF, when converting back to JSON the exponent notation is used instead.

Of the three remaining mismatches, two are caused by part of the data only being used for a join condition. The last one is a bug caused by the mapping doing a join between entities of the same type. This is something we failed to take into account when making the implementation.

4.3.2 Performance

The same testing setup as the LUBM4OBDA benchmark is used for the GTFS-Madrid-Bench. As such we observe the same limitation on the total performance due to the single threaded nature of the triple store. The results can be found in Table 4.2. We find that the time to invert scales linearly with the scale of the benchmark.

Scale	CSV	JSON
1	6.75s	12.08s
10	60.49s	109.60s
100	641.00s	1126.07s

Table 4.2 Performance on the GTFS-Madrid-Bench benchmark

Chapter 5

Conclusion

In this master thesis we explored the possibility of inverting knowledge graphs back to their original raw data format using mapping rules. To this end we created a generic two-module design and made an implementation using RML mapping rules. We implemented a data retrieval module using SPARQL and two templating modules for CSV and JSON.

The data retrieval module takes the knowledge graph, either a file or an endpoint, and uses the mapping rules to retrieve the data from it. In our case it translates selected triple maps in the mapping rules to parts of a SPARQL query, adapting the approach based on the type of the triple map. The selection process aims to balance performance and reliability. Generating the output is made more challenging by the requirement to handle both absent fields and multiple occurrences of the same field while ensuring their equality. We found that to ensure proper inversion we cannot use the full RML specification, but must impose restrictions on the mapping rules. These restrictions pertain to the access of the source and transformation of the data. To fully reconstruct a data source no aggregation, filtering or other irreversible transformations on the source can be done. Additionally, the mapping must use all the data in the source. In order to be able to reconstruct values of the data, no possibly irreversible transformations can be done on them, as any information lost during a lossy transformation is impossible to regain. This includes the use of templates, where the divider between different references cannot be present in the references. URI encoded templates are an exception, here the transformation is guaranteed to be reversible if reserved characters are used as divider.

The templating module takes the data retrieved by the data retrieval module and uses it to create the source file. In case of the CSV templating module this is a simple process, as the data is already in the correct format. The JSON templating module is more complex, as JSON has a nested structure. We take all the paths from which data is retrieved and create a template from them. This template is then recursively filled from the root node, splitting the data at each array. To properly create the JSON template all filepaths must be linearly connected to the root node, using JSON features like recursive descent makes the path taken to the data unclear and the template impossible to create. This is another restriction we must impose on the mapping rules.

We evaluated our implementation using the RML test cases to test various edge cases, the LUBM4OBDA

benchmark to test the scalability of our data retrieval module and the GTFS-Madrid benchmark to evaluate the performance of the JSON templating module. We find that our implementation is able to invert the knowledge graph back to the original source files in most cases. The failures are mostly due to limitations of the mapping processor, incompatible mapping rules or different data representations. The LUBM4OBDA benchmark shows that our data retrieval module scales linearly with the size of the knowledge graph and is completely dependent on the power of the SPARQL endpoint. The GTFS-Madrid benchmark shows that the JSON templating module also scales linearly with the size of the knowledge graph.

5.1 Future work

In this section we will discuss possible future work that can be done to expand on our implementation. The implementation can be expanded on in both depth and breadth.

5.1.1 Functions

RML offers the capability to transform data using functions with the Function Ontology (FnO). Expanding the capabilities of the data retrieval module to encompass this is a possible next step. Functions that merely transform data without losing information could be directly inverted. The data lost in lossy transformations could in some cases be regained by calling on an external data source. Lastly lossy functions could be used to verify the correctness of data values with duplicate references of which one was run through a lossy function.

Implementing this would expand the depth of the implementation. We did not implement it due to the large complexity and time to implement it compared to the limited benefits it is expected to provide.

5.1.2 Additional templating modules

We created two templating modules, one for CSV and one for JSON. Each a representative of a different type of data format: tabular and nested. Creating additional templating modules would make the implementation more versatile, further allowing mixed data formats in mappings.

A special case where we don't recreate the data source would be to make a database update module. If the primary key of a table is present in the references an update query could be made to update the database.

Bibliography

Aranda, C. B., Corby, O., and Das, S. (2013). SPARQL 1.1 overview. W3C recommendation, W3C. https://www.w3.org/TR/2013/REC-sparql11-overview-20130321/.

Arenas-Guerrero, J., Chaves-Fraga, D., Toledo, J., Pérez, M. S., and Corcho, O. (2022). Morph-KGC: Scalable knowledge graph materialization with mapping partitions. *Semantic Web*.

Arenas-Guerrero, J., Pérez, M. S., and Corcho, O. (2024). Lubm4obda: Benchmarking obda systems with inference and meta knowledge. *Journal of Web Engineering*, 22(08):1163–1186.

Bergman, M. K. (2019). A common sense view of knowledge graphs. *AI3:::Adaptive Information*, 1(1):1–1.

Berners-Lee, T., Fielding, R. T., and Masinter, L. M. (2005). Uniform Resource Identifier (URI): Generic Syntax. RFC 3986.

Berners-Lee, T., Hendler, J., and Lassila, O. (2001). The semantic web. *Scientific American*, 284(5).

Bischof, S., Decker, S., Krennwallner, T., Lopes, N., and Polleres, A. (2012). Mapping between rdf and xml with xsparql. *Journal on Data Semantics*, 1(3):147–185.

Calvanese, D., Cogrel, B., Komla-Ebri, S., Kontchakov, R., Lanti, D., Rezk, M., Rodriguez-Muro, M., and Xiao, G. (2017). Ontop: Answering sparql queries over relational databases. *Semantic Web*, 8:471–487. 3.

Chortaras, A. and Stamou, G. (2018). D2rml: Integrating heterogeneous data and web services into custom rdf graphs. In *LDOW@WWW*.

Cyganiak, R., Sundara, S., and Das, S. (2012). R2RML: RDB to RDF mapping language. W3C recommendation, W3C. https://www.w3.org/TR/2012/REC-r2rml-20120927/.

Das, S., Cyganiak, R., and Sundara, S. (2012). R2RML: RDB to RDF mapping language. W3C recommendation, W3C. https://www.w3.org/TR/2012/REC-r2rml-20120927/.

De Giacomo, G., Lembo, D., Oriol, X., Savo, D. F., and Teniente, E. (2017). Practical update management in ontology-based data access. In d'Amato, C., Fernandez, M., Tamma, V., Lecue, F., Cudré-Mauroux, P., Sequeda, J., Lange, C., and Heflin, J., editors, *The Semantic Web – ISWC 2017*, pages 225–242, Cham. Springer International Publishing.

Dimou, A., Vander Sande, M., Colpaert, P., Verborgh, R., Mannens, E., and Van de Walle, R. (2014). RML: a generic language for integrated RDF mappings of heterogeneous data. In Bizer, C., Heath, T., Auer, S., and Berners-Lee, T., editors, *Proceedings of the 7th Workshop on Linked Data on the Web*, volume 1184 of *CEUR Workshop Proceedings*.

Guo, Y., Pan, Z., and Heflin, J. (2005). Lubm: A benchmark for owl knowledge base systems. *J. Web Semant.*, 3(2-3):158–182.

Harris, S. and Seaborne, A. (2013). SPARQL 1.1 query language. W3C recommendation, W3C. https://www.w3.org/TR/2013/REC-sparql11-query-20130321/.

Heyvaert, P., De Meester, B., Dimou, A., and Verborgh, R. (2018a). Declarative rules for linked data generation at your fingertips! In Gangemi, A., Gentile, A. L., Nuzzolese, A. G., Rudolph, S., Maleshkova, M., Paulheim, H., Pan, J. Z., and Alam, M., editors, *The Semantic Web: ESWC 2018 Satellite Events*, pages 213–217, Cham. Springer International Publishing.

Heyvaert, P., Dimou, A., De Meester, B., Seymoens, T., Herregodts, A.-L., Verborgh, R., Schuurman, D., and Mannens, E. (2018b). Specification and implementation of mapping rule visualization and editing: MapVOWL and the RMLEditor. *Journal of Web Semantics*, 49:31–50.

Heyvaert, P., Dimou, A., and Meester, B. D. (2019). Rml test cases. `https://github.com/kg-construct/rml-test-cases`. Accessed: 2023-12-21.

Ji, S., Pan, S., Cambria, E., Marttinen, P., and Yu, P. S. (2022). A survey on knowledge graphs: Representation, acquisition, and applications. *IEEE Transactions on Neural Networks and Learning Systems*, 33(2):494–514.

Manola, F. and Miller, E. (2004). RDF primer. W3C recommendation, W3C. https://www.w3.org/TR/2004/REC-rdf-primer-20040210/.

Marchi, E. and Miguel, O. (1974). On the structure of the teaching-learning interactive process. *International Journal of Game Theory*, 3(2):83–99.

Meester, B. D., Heyvaert, P., and Delva, T. (2022). RML: RDF mapping language. Unofficial draft, RML. https://rml.io/specs/rml/.

Michel, F., Djimenou, L., Faron Zucker, C., and Montagnat, J. (2015). Translation of relational and non-relational databases into rdf with xr2rml. In *WEBIST (Selected Papers)*, pages 443–454.

Polleres, A., Lopes, N., Decker, S., Krennwallner, T., and Kopecky, J. (2009). Xsparql language specification. Technical report, W3C. https://www.w3.org/submissions/xsparql-language-specification/.

Priyatna, F., Chaves, D., Toledo, J., Doña, D., Ruckhaus, E., Assche, D. V., and Guerrero, J. A. (2023). oeg-upm/gtfs-bench: v1.3.0.

Richens, R. H. (1956). Preprogramming for mechanical translation. *Mech. Transl. Comput. Linguistics*, 3:20–25.

Seaborne, A. and Prud'hommeaux, E. (2008). SPARQL query language for RDF. W3C recommendation, W3C. https://www.w3.org/TR/2008/REC-rdf-sparql-query-20080115/.

Singhal, A. (2012). Introducing the knowledge graph: things, not strings. 2020-11-13.

Spieldenner., D. (2022). Poser: A semantic payload lowering service. In *Proceedings of the 18th International Conference on Web Information Systems and Technologies - WEBIST*, pages 249–256. INSTICC, SciTePress.

Unbehauen, J. and Martin, M. (2017). Sparql Update queries over R2RML mapped data sources. In Eibl, M. and Gaedke, M., editors, *INFORMATIK 2017, Lecture Notes in Informatics (LNI)*, Lecture Notes in Informatics (LNI), pages 1891–1901, Chemnitz, Germany. Gesellschaft für Informatik.

Van Assche, D., Delva, T., Haesendonck, G., Heyvaert, P., De Meester, B., and Dimou, A. (2023). Declarative rdf graph generation from heterogeneous (semi-)structured data: A systematic literature review. *Journal of Web Semantics*, 75:100753.

Vu, B., Pujara, J., and Knoblock, C. A. (2019). D-repr: A language for describing and mapping diversely-structured data sources to rdf. In *Proceedings of the 10th International Conference on Knowledge Capture*, K-CAP '19, pages 189–196, New York, NY, USA. Association for Computing Machinery.